much. CONSULTING

SIMULATION REPORT — ROUND 2

Post-Revision Audit & Simulation

Odoo ERP Architecture & Pitch Package — MDR Devices Medizintechnik GmbH

HIRING ENTITY
much. Consulting

BASIS
Independent re-audit against the current state of all ten package documents, run fresh rather than extended from Round 1's findings — Round 1 is checked for whether its own action items actually landed, not assumed.

1 · ROUND 1 → ROUND 2: WHAT ACTUALLY GOT FIXED

Round 1 left four open action items. Verified against the current document state:

ROUND 1 ACTION ITEM	STATUS	WHERE
Rewrite the Slide 12 / memo close to a committed ask	DONE	Slide 12 (blueprint, script, and the actual `.pptx`) now states fixed-fee/3-week scope and a decision trigger. Target date and fee are intentionally left as `[bracketed placeholders]` — see Section 6, this is the one piece only MDR Devices' actual timeline can fill in.
Explicitly place multi-tier payer reconciliation in the extension layer, not just in the rebuttal playbook	DONE	Object Schema §4.2, the Strategic Memo's settlement bullet, and the GxP plan's Category 5 list and regression suite (TC-12) now all say this consistently. Previously only the rebuttal playbook admitted it — the exact inconsistency flagged in the prior session.
Add reverse consignment returns, reusable-instrument sterilization, and legacy cutover to the documents	DONE restructured rather than pasted in as-is	Reverse returns and the instrument loop are now Stage 4B and a Parallel Track in the Value Stream doc, with matching fields in Object Schema §5. Cutover was moved to the GxP Validation Plan (§6) instead of the Value Stream doc — cutover is a one-time project event, not a steady-state process stage, and burying it in the operational flow doc was the wrong home for it even before content was added.
Correct the Quality Hold field mapping and add the fail-after-hold disposition path	DONE	Object Schema §3.3 already had Reordering Rules correctly separated from the hold mechanism; the missing piece — disposition when a Quality Check fails, before or after the 90-day timer — is now a table row, with a corresponding regression test (TC-13).

All four are real, verifiable changes in the live documents, not just claims — each is cited to its exact location above so it can be checked directly.

2 · NEW FINDINGS THIS PASS

Fresh eyes on the *current* state, not a re-run of Round 1's list:

2.1 — ROUND 1'S OWN "FIX" TO THE CONSIGNMENT MECHANISM WAS ITSELF WRONG, AND HAS BEEN CORRECTED

Round 1 (Finding 3) redirected the package to lead with Odoo's `owner_id` field as the native consignment mechanism. Checked against Odoo's actual documentation: `owner_id`/the "Consignment" feature is built for the *inbound* direction — stock in your warehouse that a vendor still owns. MDR Devices needs the mirror image — stock it still owns, sitting at a hospital. The mechanism that actually does that job is location **type** (internal vs. customer), not ownership tagging. This is now corrected in Object Schema §1.3 and Value Stream Stage 4, with a flag to verify hands-on before presenting it, since this is a derived correction, not something tested in a live instance. This matters beyond the technical detail: it shows why "fixed by a previous pass" can't be taken as "verified" — Round 1's confidence on this specific point was misplaced, and would have been a bad moment to have stated flatly in front of a CIO who has actually run Odoo consignment.

2.2 — DEPLOYMENT TIER WAS NEVER SPECIFIED, AND THE VALIDATION CLAIMS DEPEND ON IT

Every document asserts a controlled, reviewable upgrade path. That claim is only fully true on **Odoo.sh** (or on-premise), where a staging branch can be tested against production data before a controlled promotion. **Odoo Online** applies platform upgrades automatically on Odoo's own schedule — Studio customizations still need a compatibility pass, but the timing of the underlying change isn't fully in MDR Devices' control. The package now names this explicitly (GxP plan §7) and states the assumption (Odoo.sh) rather than leaving the tier implicit. A technical CIO is likely to ask which tier is actually being proposed; the pitch should have an answer ready that isn't "we hadn't decided."

2.3 — GDPR IS A LIVE GAP, ENTIRELY SEPARATE FROM MDR/GXP, AND WAS NEVER MENTIONED ANYWHERE IN THE PACKAGE

The case names patient information and physician records living in the CRM. That's special-category health data under GDPR Article 9 — a different legal regime from MDR/GxP, with its own legal-basis, minimization, retention, and likely DPIA requirements. Ten documents, zero mentions of data protection, before this pass. Now flagged in the GxP plan (§8) and given a rebuttal entry (Objection 5C) — deliberately scoped as "raise it, route it to legal/DPO," not solved, since a fabricated GDPR design would be worse than an honest gap.

2.4 — STILL OPEN, NOT YET ACTIONED: UNMATCHED RETURNS

Stage 4B assumes a returned lot can be matched back to its `stock.lot` by UDI scan. Real-world edge case worth having an answer for: a hospital return arrives with incomplete paperwork or a damaged/unreadable label — no scan, no clean match. The current design has no fallback path for a return that can't be automatically reconciled to a lot. Worth a sentence in discovery framing ("manual match-and-flag queue, reviewed by Quality") rather than presenting the automated path as the only path. Not fixed in this pass — flagged for the next one, listed in Section 6.

3 · UPDATED BOARDROOM SIMULATION

Two of the five new objections this round are now written into the rebuttal playbook itself (Objections 5A, 5B, 5C) so they live with the rest of the prep, not just in this report:

- **5A (CIO)** — reusable instruments aren't implants, where's that in the design
- **5B (CEO)** — what happens to stock that comes back
- **5C (Partner)** — MDR/GxP is covered, what about GDPR

One more, specific to this round's own technical correction, worth being ready for since it's the kind of thing that separates "read the deck" from "understands the deck":

Q6 Testing Whether the Candidate Actually Understands the Mechanism or Is Reciting It

CIO

"You've told me hospital-held stock stays on your balance sheet. Walk me through why — what specifically in Odoo makes that true?"

MODEL ANSWER

"It's the location type, not an ownership tag. Odoo values stock based on whether it's sitting in an internal location — internal to us, regardless of physical distance — versus a customer location, which is what an actual sale moves it to. The hospital consignment location is modeled as internal, so it stays on our books the whole time it's physically there. The `owner_id` field people sometimes reach for here is actually built for the reverse situation — someone else's stock sitting in your warehouse — which isn't what's happening at MDR Devices. Getting the direction right here isn't a technicality; it's the difference between an audit-defensible valuation model and one that only happens to look right on a slide."

This question is worth preparing for specifically because it probes whether Round 1's error would have been caught live — the honest answer is "not necessarily," which is exactly why it's called out here rather than quietly folded in.

4 · UPDATED VIABILITY ASSESSMENT

~78%

PITCH PROBABILITY OF WIN (ROUND 1: 68%)

68	78	—	—
R1	R2	R3	R4

8/10

HIRING PANEL SCORE (ROUND 1: 7/10)

●●●●●●●●○○

DRIVER	ROUND 1	ROUND 2	WHY IT MOVED
Commercial close	−8%	−2%	Structurally specific now (scope, fixed-fee framing, decision trigger); residual penalty only because date/fee are still placeholders MDR Devices' real calendar has to fill in
Multi-tier payer honesty	−6%	0%	Now consistent across memo, architecture, validation plan, and rebuttal — closed, not just admitted in one document
Reverse logistics + instrument reprocessing	−10%	−2%	Both modeled with fields, flow stages, and rebuttal answers; residual penalty because none of it is hands-on tested in a live Odoo instance
Technical accuracy	+12%	+14%	Round 1's own correction (`owner_id`) has itself been corrected; net technical grounding is higher than either prior state, but see below
Newly-introduced risk: unverified technical claim	n/a	−2%	The location-vs-ownership correction is well-reasoned but derived, not sandbox-verified — flagged plainly rather than stated with false certainty
Newly-found gap: GDPR	n/a	−3%	Real, previously entirely absent; now named and routed correctly, but not designed — a thorough panel can still press on it
Deployment tier clarity	n/a	+2%	Removes a plausible, previously-unanswered CIO question
Fluff / AI-pattern language	+5%	+5%	No change — still clean
Rebuttal depth and coverage	+10%	+12%	Two new objections written into the actual playbook, not just this report, plus the meta-question in Section 3

This is a qualitative, reasoned estimate, not a computed model — the table shows the reasoning so it can be argued with, not to dress up a guess as precision.

Hiring Panel Score: 8/10 for a Principal Architect seat (Round 1: 7/10). What separated 7 from 8 last round — catching reverse logistics and the instrument loop without being told — is now closed. What separates 8 from 9–10 isn't on paper: live command of the room under the Q6-style question above, and the fact that nothing in this package has touched a real Odoo instance yet. A Principal-level hire is expected to know the difference between "this is well-reasoned" and "this is verified" — this report tries to model that distinction rather than paper over it.

5 · UPDATED COMPLIANCE & REASONING MATRIX

CORE REQUIREMENT	ARCHITECTURE / FIELD MAPPING	REGULATORY VIABILITY	REAL-WORLD EDGE CASE
MDR/UDI Traceability	`stock.lot`, native GS1 parsing + extension-layer DI/PI sub-fields	**PASS** conditional on extension delivery	Hospitals relabeling with proprietary internal tags that don't conform to GS1 — no documented fallback for a non-parsing scan. Not yet fixed.
90-Day Quality Hold	Quality Control Point + push rule + scheduled action; failure disposition now modeled	**PASS**	Closed this round (Failure Disposition row, TC-13).
Consignment Valuation	Internal-location type (corrected this round from an incorrect `owner_id`-led framing)	**PASS** pending hands-on sandbox confirmation	The one claim in this package derived rather than verified — flag before stating it with full confidence live.
Physician Gatekeeping	`res.partner` credential fields + `sale.action_confirm()` override	**PASS**	Locum/visiting physicians without an existing record — needs a "pending verification" state, not just validated/expired/suspended. Not yet fixed.
Multi-Stage Claims	`account.move` native two-way split + named extension for 3+ parties/rebate netting	**PASS**	Closed this round across all documents (previously only the rebuttal playbook said this).
Reusable Instrument Lifecycle	New: Object Schema §5, Value Stream Parallel Track	**PASS** conditional on extension delivery	Cross-hospital tray sharing (a tray used at Hospital A one week, Hospital B the next) isn't addressed — worth confirming during discovery whether that occurs.
Data Protection (GDPR)	Not yet designed — flagged and routed to legal/DPO	**AT RISK** newly identified	Entirely new this round. See Section 2.3.

6 · WHAT'S STILL GENUINELY OPEN

Not closed by this pass, and shouldn't be presented as if they are:

- [ ] Fill in the actual target decision date and fee structure on Slide 12 (blueprint, script, and `.pptx`) — the only remaining piece only MDR Devices' real calendar and much. Consulting's real commercial terms can supply.
- [ ] Sandbox-verify the internal-location consignment valuation claim (Section 2.1 / Q6) before stating it with full confidence live.
- [ ] Decide and confirm the actual deployment tier (Odoo.sh vs. Online vs. on-premise) being proposed — currently an assumption, not a decision.
- [ ] Unmatched-return fallback path (Section 2.4) — not yet designed, only flagged.
- [ ] Locum/visiting physician "pending verification" state (compliance matrix, Physician Gatekeeping row) — carried over from Round 1, still open.
- [ ] One manual visual QA pass on the `.pptx` in actual PowerPoint/Keynote — Slides 5 and 12 both grew new text this round; a rendered check confirms nothing wraps awkwardly at presentation resolution, on top of the automated render check already run against this

edit.

- [ ] Rehearse the two new boardroom objections (5A, 5B, 5C) and the Q6 mechanism question out loud, at full pace — none of this round's writing has been spoken yet.

7 · NET ASSESSMENT

Round 1 correctly identified real gaps and made real improvements, but also introduced one confident-sounding error of its own (Section 2.1) — a useful reminder that “audited” isn't the same as “correct,” which is the reason this pass was run independently rather than as an extension of the last one. What's in the package now is measurably stronger: two structural gaps closed with actual field-level design, one technical error caught and fixed, one previously-invisible compliance dimension named, and the rebuttal playbook grown to match. What's left is exactly what document editing can't finish — a real number on the ask, a hands-on check of the one derived technical claim, and rehearsal. That's a good place to be with 24 hours left, not a reason to stop here.