

Cross-Reference Audit: External Package Review

Systematic comparison against an independently-generated pitch package for the same case, produced in a different AI tool

METHOD

Every file in the external package read in full, checked claim-by-claim against this package and, where load-bearing, against outside verification — not adopted on the strength of “another model spent a lot of time on it.”

1 • PROVENANCE, ESTABLISHED FIRST

Both packages open with byte-for-byte identical 06_Audit_and_QA and 07/01_Report documents. They share a common Round-1 origin and diverged from there. That explains the structural similarity and, more importantly, explains why one specific error appears in both lineages — it was introduced at the shared origin point and only caught in one of the two branches. Worth knowing before assuming convergence between two AI outputs means correctness; here, convergence on the same wrong answer had a common cause, not independent verification.

The external package also carries its own internal history: a rejected “Repository B” (a third, separately-sourced document containing self-graded claims — a fabricated “15–20 person days to under 4 hours” efficiency statistic, and referenced document stages that didn’t actually exist in the package). The external package’s own audit trail caught this and explicitly declined to incorporate it. That’s a correct call, made independently of this review, and it’s noted here only because it’s relevant to Section 3 below.

2 • IMPORTED — VERIFIED USEFUL, NOW IN THIS PACKAGE

ADDITION	WHAT IT CLOSES	WHERE IT LANDED
Pending Verification credential state for locum/visiting physicians	A gap this package's own Round 2 report flagged as open — now actually fixed, not just named	Object Schema §1.2
Manual-entry UDI fallback for non-GS1 hospital relabeling	Another gap Round 2 flagged as open, now fixed, with an x_udi_source marker added on top of the external version to keep scanned vs. hand-entered records distinguishable in reporting	Object Schema §2.4
Explicit x_return_disposition / x_return_reason fields	Made the reverse-logistics flow schema-level and queryable rather than purely narrative	Object Schema §1.3, Value Stream Stage 4B, two new regression tests
Autoclave cycle parameters (temperature, duration, pressure) captured on the quality check	More credible than a bare pass/fail flag on a sterilization check	Object Schema §5.3, Value Stream Parallel Track
Margin-recovery framing on vendor rebates (“recovery, not projection”)	Sharper, more specifically defensible CFO argument than a generic efficiency claim	Object Schema §4.3, Strategic Memo
QA headcount redeployment framing (evidence-assembly hours → CAPA prevention)	Concrete answer to “what do the recovered hours actually do”	Strategic Memo
“MedTech Compliance Accelerator” — practice-asset framing tied to fee justification	Reframes the fee as buying a refined pattern, not a one-off build; a genuinely good Partner-facing angle	Strategic Memo close, Slide 12 (blueprint, script, and the actual .pptx — added as a new, separately verified text element)

ADDITION	WHAT IT CLOSES	WHERE IT LANDED
Signature line: “We don't sell zero custom code. We sell named custom code.”	A memorable, quotable summary of the whole native-core/named-extension positioning	Slide 5 (spoken emphasis — deliberately not printed on the slide)
Default parallel-run window (shorter of 90 days or oldest open record) + explicit read-only-archive status for the legacy system post-cutover	Makes the cutover plan concrete instead of “discovery will decide,” while still framed as a proposed default, not a fabricated fact about MDR Devices	GxP Validation Plan §6
Four new rebuttal objections: sterilization data-integrity, return/scrap cost allocation (system vs. contract), cutover cross-reference for new activity against a migrated device, and — the sharpest one — being asked to name and let someone call a referenced “comparable engagement”	The reference-client question is arguably the single best new item from this whole review: neither package had it before, and an unprepared answer here is a genuine pitch-killer	Rebuttal Playbook, Objections 5D-5G

3 • REVIEWED AND REJECTED — WITH REASONS

3.1 — THE `owner_id` -AS-PRIMARY-MECHANISM ERROR

The external package still states, in its current live architecture document, that Odoo's `owner_id` field is the primary native mechanism for keeping hospital-held consignment stock on MDR Devices' balance sheet. This is the same error this package caught and corrected in Round 2 (`owner_id` is documented for the *inbound* direction — stock you receive that a vendor still owns; MDR Devices needs the outbound/consignor direction, which internal-location type handles). It persisted across all of the external package's rounds, including its most recent one. **Not imported.** This is worth knowing on its own merits: a second, differently-produced package independently arrived at the same plausible-sounding but incorrect answer, which is a mild point of caution about how confidently-worded AI output can be wrong in ways that don't announce themselves.

3.2 — A CLAIMED “1-PAGE MEMORANDUM” REQUIREMENT

The external package's own later-round audit asserts the original case specification calls for a one-page pre-read memorandum, and uses that to justify trimming its memo. The actual case document (re-checked directly for this review) contains no page-count, word-count, or length constraint of any kind — only “send over prepared documents” and a 30-minute pitch slot. **Not imported as a constraint.** This package's memo is being kept concise on its own merits — it's 749 words after this round's additions, well short of the external package's own reported 979 — but not because of a rule that doesn't exist in the source material. Stating a fabricated constraint as fact would have been a real risk if repeated in front of the panel that actually wrote the case.

3.3 — REPOSITORY B'S FABRICATED STATISTICS

Already rejected upstream, in the external package's own audit trail. Confirmed here as correctly excluded, and not reintroduced from either source. Included in this report only for completeness — it's the same discipline being applied twice.

4 • FOUND ALONG THE WAY — NOT FROM THE EXTERNAL PACKAGE

Cross-referencing forces a full re-read of this package's own files, which surfaced two of this package's own Round 2 inconsistencies that had nothing to do with the external review: the Slide 5 blueprint and verbatim script still described the extension layer as three items after the actual `.pptx` had already been updated to five. Both are now fixed and consistent across blueprint, script, and deck. Worth naming plainly: this round improved this package's *internal* consistency as much as it imported anything external.

5 • A NOTE WORTH REHEARSING, NOT JUST READING

Being able to describe this exact process to the Partner persona — checked a second AI-produced attempt at the same case, found real gaps worth closing, caught a real technical error it repeated across four rounds, and caught a fabricated constraint in its own audit trail — is itself a strong, specific answer to “how do you use AI in your work,” if that question comes up. It’s a concrete demonstration of exactly the judgment a Principal Architect role is testing for: not “did you use the tool,” but “did you verify what it gave you.” Worth having a tight, honest version of this ready, not as a scripted brag, but because it’s true and it directly answers the question a skeptical Partner is actually asking.

6 • UPDATED ASSESSMENT



A real but bounded move, not a re-rating from scratch. What improved: two previously-open gaps are now genuinely closed, the commercial argument is sharper and more specific, the rebuttal playbook now covers the single sharpest question either package had missed, and the package correctly avoided importing a technical error and a fabricated constraint that a less careful pass might have absorbed. What didn't move at all, and is the ceiling on how far this number can honestly go: nothing in either package has touched a live Odoo instance, and none of this — including everything added this round — has been rehearsed aloud.

Hiring Panel Score deliberately unchanged. Round 2 named the specific reason this wasn't a 9: live command of the room under pressure, and zero hands-on verification. Nothing in this round addressed either of those — this was a content and cross-referencing pass, not a rehearsal or a sandbox session. Raising the score without addressing the stated reason for the cap would be exactly the kind of ungrounded number inflation this whole review process has been checking for elsewhere. The content is stronger; the reason for the ceiling hasn't moved.

7 • STILL OPEN

Unchanged in kind from Round 2, and worth repeating rather than letting a longer document imply they've been resolved:

- ☐ Real date and fee on Slide 12 — still bracketed
- ☐ Sandbox-verify the consignment valuation mechanism and the new return-disposition field
- ☐ GDPR — named and routed, still not designed
- ☐ Confirm deployment tier (Odoo.sh vs. Online) being proposed
- ☐ Rehearse — and this round added more new material (the signature line, the accelerator framing, four new objections) than any single prior round. More has been written; none of it has been said out loud yet.