

The Whole Thing, Assessed

Deliverable, live system and pitch — measured against the case brief, with an honest scorecard and the residual weak points named

Prepared by Waqas Abbas Mian · Senior ERP Solution Architect — final assessment before pitch

1 · What was assessed, and how

This is the top-level verdict across everything: the fifteen-document client deliverable, the live Odoo demonstration, and the pitch material — measured against the actual case brief, the running system, and the applicable regulation. It draws on a live regression pass (18/18, executed against the instance), a word-by-word peer review, a forensic text/visual sweep, and a brief-coverage decode. No score below is asserted without a source.

2 · The scorecard

DIMENSION	RATING	BASIS
Brief coverage	21 / 21	Every explicit and implied requirement in the brief maps to a deliverable and, where demonstrable, a live proof — verified by text search across all fifteen documents. The RTM is the built-in coverage map.
Live system integrity	18 / 18	Regression pass executed against the instance: four compliance gates (code-verified + state-consistent), UDI, genealogy, consignment valuation, analytic spine, reorder scoping, returns, audit trail, locations, quality — all pass; demo left pristine.
Solution viability	High	Native-where-covered, five named extensions where not; mixed GAMP classification defended; phased, fixed-fee-discovery commercial shape with a statutory-deadline Horizon 2 tail. No claim exceeds what the live system does.
Text & language quality	Clean	Zero mojibake, zero AI attribution, uniform British English across the client deliverable, all metadata authored correctly — confirmed by forensic sweep and word-by-word peer read.
Visual & layout quality	Clean	Consistent house style; no dead white space or truncation; diagrams (process map, architecture, roadmap) render as one-page exhibits; page-by-page render inspection passed.
Honesty / defensibility	Distinctive	Every claim labelled demonstrated / design-ready / discovery-scoped; every platform boundary declared (e-signature, Field Service, payer-split, valuation wiring). The teardown simulation surfaced no undefended claim.
Differentiation	Strong	A live built system (not slides), an AI working-method note with a correction log, and the UK Compliance Bridge mapping the thread onto the 2025–26 UK device reforms — research no other candidate is likely to bring.

3 · Why the solution is viable, in one pass

The proposal answers the brief's real question — "what does one compliant end-to-end process look like?" — with a single validated data thread across six departments, proven on a running system rather than asserted on a slide. It is commercially coherent: fixed-fee discovery de-risks entry, value lands inside the first configuration sprint, and Horizon 2 attaches a revenue tail to a statutory deadline the client cannot avoid. Crucially, after the live regression and the earlier

document-versus-system reconciliation, **there is no gap between what the package says and what the instance does** — which removes the one failure mode that sinks a live pitch: a claim the demo contradicts.

4 · The residual weak points — named, not buried

WEAK POINT	HONEST STATUS & MITIGATION
Demo runs on a trial tenant	Can expire or exhibit tenant quirks (the API-confirm compute issue is one). Mitigated by the double pre-flight check and the UAT guide as a screenshot fallback. It is a tenant characteristic, not a solution defect.
Field Service simulated (T-14)	The one requirement the brief names twice, unavailable on the tier — declared SIMULATED, with the FSM production design stated. Own it before it is asked.
Fee is a placeholder	Deliberate — a number before discovery is a guess dressed as a promise. Have the reasoning ready for the inevitable "ballpark?".
GDPR / patient-data depth	Special-category data named as a distinct legal workstream for the DPO — shape given, not solved. Correct posture; do not overclaim.
Two stray draft invoices	Downstream of a confirmed order, unposted, touching no compliance claim — noted in the regression report. A 30-second UI tidy before the demo if desired.
German localisation / PLM (R&D)	Study-level, not delivery-level, and scoped rather than built — the honest-gap answers are prepared in the Q&A drill.

5 · Verdict

READINESS

The deliverable is **submission-ready and pitch-ready**. It covers the brief completely, it is true on the running system, it reads and looks clean, and it survives an adversarial teardown because its honesty is engineered in. On the panel's own axes — competence, client-presence, judgement and fit — the package sits in the **upper band**, and the residual risks are all either declared boundaries or one-line answers already rehearsed. The remaining variable is not the artefact; it is the delivery — the out-loud reps against the clock.

THE ONE THING TO REMEMBER

Nothing in this package needs to be oversold, because everything in it is either demonstrated or honestly scoped. Walk in and let the system — and the honesty — do the work.